

# Akhil\_Vallala\_Project\_1

March 1, 2024

Akhil Vallala's Project 1

## 1 Importing packages

```
[4]: # Array math
import numpy as np

# Iteration tracking
from tqdm import tqdm

# Type hinting
from typing import Tuple

# Plotting
import matplotlib.pyplot as plt
import matplotlib
import seaborn as sns

# Os traversal
import os
```

For rendering Greek letters in Latex for plotting

```
[5]: matplotlib.rcParams['font.family'] = 'sans-serif'
matplotlib.rcParams['text.usetex'] = True
matplotlib.rcParams["text.latex.preamble"] = r"\usepackage{cmbright}"
matplotlib.rcParams.update(matplotlib.rcParamsDefault)
```

## 2 Defining the plotting functions

### 2.1 Array Index to Matplot Coords Converter

```
[6]: def array_index_to_matplot_coords(i: int, j: int, n_cols: int) -> Tuple[int,
↪int]:
    """Converts an array index to a matplot coordinate"""
    x = j
    y = n_cols - i - 1
```

```
return x, y
```

## 2.2 Maze Plotter

```
[7]: def plot_matrix(
    M: np.array,
    goal_coords: list = [],
    hole_coords: list = [],
    start_coords: list = [],
    highlight_coords: list = [],
    img_width: int = 5,
    img_height: int = 5,
    title: str = None,
    filename: str = None,
) -> None:
    """
    Plots a matrix as an image.
    """
    height, width = M.shape

    fig = plt.figure(figsize=(img_width, img_width))
    ax = fig.add_subplot(111, aspect='equal')

    for x in range(height):
        for y in range(width):
            # By default, the (0, 0) coordinate in matplotlib is the bottom
            ↪left corner,
            # so we need to invert the y coordinate to plot the matrix correctly
            matplot_x, matplot_y = array_index_to_matplot_coords(x, y, height)

            # If there is a tuple of (x, y) in the goal_coords list, we color
            ↪the cell gray
            if (x, y) in goal_coords:
                ax.add_patch(matplotlib.patches.Rectangle((matplot_x - 0.5,
            ↪matplot_y - 0.5), 1, 1, facecolor='lightgreen'))
            # If there is a tuple of (x, y) in the hole_coords list, we color
            ↪the cell salmon
            elif (x, y) in hole_coords:
                ax.add_patch(matplotlib.patches.Rectangle((matplot_x - 0.5,
            ↪matplot_y - 0.5), 1, 1, facecolor='salmon'))
            # If there is a tuple of (x, y) in the start_coords list, we color
            ↪the cell yellow
            elif (x, y) in start_coords:
                ax.add_patch(matplotlib.patches.Rectangle((matplot_x - 0.5,
            ↪matplot_y - 0.5), 1, 1, facecolor='yellow'))
```

```

        # If there is a tuple of (x, y) in the highlight_coords list, we
        ↪color the cell lightblue
        elif (x, y) in highlight_coords:
            ax.add_patch(matplotlib.patches.Rectangle((matplot_x - 0.5,
            ↪matplot_y - 0.5), 1, 1, facecolor='lightblue'))

            ax.annotate(str(M[x][y]), xy=(matplot_x, matplot_y), ha='center',
            ↪va='center')

    offset = .5
    ax.set_xlim(-offset, width - offset)
    ax.set_ylim(-offset, height - offset)

    ax.hlines(y=np.arange(height+1)- offset, xmin=-offset, xmax=width-offset)
    ax.vlines(x=np.arange(width+1) - offset, ymin=-offset, ymax=height-offset)

    plt.title(title)
    if filename is not None:
        plt.savefig(filename)
        plt.close(fig)
    else:
        plt.show()

```

```

[8]: def plot_state_visits(
    S: np.array,
    visit_dict: dict,
    img_width: int = 5,
    img_height: int = 5,
    ):
    """
    Plots the states and colors them by the number of visits.

    The more visits a state has, the darker the color.
    """
    # Get the number of rows and columns
    n_rows, n_cols = S.shape
    # Create a new matrix to hold the number of visits
    V = np.zeros((n_rows, n_cols))
    # Iterate through the visit dictionary and update the V matrix
    for s, visits in visit_dict.items():
        # Converting the state to an array index
        s_index = np.where(S == s)

        row, col = s_index[0][0], s_index[1][0]

        V[row, col] = visits

```

```

fig = plt.figure(figsize=(img_width, img_height))
ax = fig.add_subplot(111, aspect='equal')

# Plotting the matrix
sns.heatmap(V, cmap='Blues', cbar=False, annot=True, fmt='.0f', ax=ax)
plt.show()

```

## 2.3 Policy Plotter

```

[9]: # Defining the plotting function for the policy
def plot_policy(
    S: np.ndarray,
    policy: dict,
    goal_coords: tuple,
    wall_coords: list,
    start_coords: tuple,
    optimal_policy_path: list,
    title: str,
    img_width: int = 6,
    img_height: int = 6
):
    height, width = S.shape

    fig = plt.figure(figsize=(img_width, img_height))
    ax = fig.add_subplot(111, aspect='equal')
    for x in range(height):
        for y in range(width):
            matplot_x, matplot_y = array_index_to_matplot_coords(x, y, height)

            # If there is a tuple of (x, y) in the hole_coords list, we color
            ↪ the cell gray
            if (x, y) in wall_coords:
                ax.add_patch(matplotlib.patches.Rectangle((matplot_x - 0.5,
            ↪ matplot_y - 0.5), 1, 1, facecolor='gray'))

            # If there is a tuple of (x, y) in the goal_coords list, we color
            ↪ the cell yellow
            elif (x, y) in goal_coords:
                ax.add_patch(matplotlib.patches.Rectangle((matplot_x - 0.5,
            ↪ matplot_y - 0.5), 1, 1, facecolor='lightgreen'))

            else:
                try:
                    # If there is a tuple of (x, y) in the start_coords list,
                    ↪ we color the cell green
                    if (x, y) in start_coords:

```

```

        ax.add_patch(matplotlib.patches.Rectangle((matplot_x -
↪0.5, matplot_y - 0.5), 1, 1, facecolor='yellow'))

        # Adding the arrows to the plot
        if 0 == policy[S[x, y]]:
            plt.arrow(matplot_x, matplot_y, 0, 0.3, head_width
↪= 0.05, head_length = 0.05)
        if 1 == policy[S[x, y]]:
            plt.arrow(matplot_x, matplot_y, 0, -0.3, head_width
↪= 0.05, head_length = 0.05)
        if 2 == policy[S[x, y]]:
            plt.arrow(matplot_x, matplot_y, -0.3, 0, head_width
↪= 0.05, head_length = 0.05)
        if 3 == policy[S[x, y]]:
            plt.arrow(matplot_x, matplot_y, 0.3, 0, head_width
↪= 0.05, head_length = 0.05)

    elif (x, y) in optimal_policy_path:
        ax.add_patch(matplotlib.patches.Rectangle((matplot_x -
↪0.5, matplot_y - 0.5), 1, 1, facecolor='lightblue'))

        # Adding the arrows to the plot
        if 0 == policy[S[x, y]]:
            plt.arrow(matplot_x, matplot_y, 0, 0.3, head_width
↪= 0.05, head_length = 0.05)
        if 1 == policy[S[x, y]]:
            plt.arrow(matplot_x, matplot_y, 0, -0.3, head_width
↪= 0.05, head_length = 0.05)
        if 2 == policy[S[x, y]]:
            plt.arrow(matplot_x, matplot_y, -0.3, 0, head_width
↪= 0.05, head_length = 0.05)
        if 3 == policy[S[x, y]]:
            plt.arrow(matplot_x, matplot_y, 0.3, 0, head_width
↪= 0.05, head_length = 0.05)
    except Exception as e:
        print(f"Error: {e}")
        print(f"Current x and y: {x}, {y}")

offset = .5
ax.set_xlim(-offset, width - offset)
ax.set_ylim(-offset, height - offset)

ax.hlines(y=np.arange(height+1)- offset, xmin=-offset, xmax=width-offset)
ax.vlines(x=np.arange(width+1) - offset, ymin=-offset, ymax=height-offset)

plt.title(title)

```

```
plt.show()
```

## 2.4 Q Table Plotter

```
[10]: # Defining the plotting function for the policy
def plot_Q_Table(
    S: np.ndarray,
    Q: np.ndarray,
    goal_coords: tuple,
    wall_coords: list,
    start_coords: tuple,
    title: str,
    img_width: int = 6,
    img_height: int = 6,
    filename: str = None
):
    height, width = S.shape

    fig = plt.figure(figsize=(img_width, img_height))
    ax = fig.add_subplot(111, aspect='equal')
    for x in range(height):
        for y in range(width):
            matplot_x, matplot_y = array_index_to_matplot_coords(x, y, height)

            # If there is a tuple of (x, y) in the hole_coords list, we color
            ↪ the cell gray
            if (x, y) in wall_coords:
                ax.add_patch(matplotlib.patches.Rectangle((matplot_x - 0.5,
            ↪ matplot_y - 0.5), 1, 1, facecolor='gray'))

            # If there is a tuple of (x, y) in the goal_coords list, we color
            ↪ the cell yellow
            elif (x, y) in goal_coords:
                ax.add_patch(matplotlib.patches.Rectangle((matplot_x - 0.5,
            ↪ matplot_y - 0.5), 1, 1, facecolor='lightgreen'))

            else:
                try:
                    # If there is a tuple of (x, y) in the start_coords list,
                    ↪ we color the cell green
                    if (x, y) in start_coords:
                        ax.add_patch(matplotlib.patches.Rectangle((matplot_x -
            ↪ 0.5, matplot_y - 0.5), 1, 1, facecolor='yellow'))
```

```

plt.arrow(matplot_x, matplot_y, 0, 0.2, head_width = 0.05,
↪head_length = 0.05)
ax.annotate(f"{Q[S[x, y]][0]:.2f}", xy=(matplot_x,
↪matplot_y + 0.5), ha='center', va='top')

plt.arrow(matplot_x, matplot_y, 0, -0.2, head_width = 0.05,
↪head_length = 0.05)
ax.annotate(f"{Q[S[x, y]][1]:.2f}", xy=(matplot_x,
↪matplot_y - 0.5), ha='center', va='bottom')

plt.arrow(matplot_x, matplot_y, -0.2, 0, head_width = 0.05,
↪head_length = 0.05)
ax.annotate(f"{Q[S[x, y]][2]:.2f}", xy=(matplot_x - 0.5,
↪matplot_y), ha='left', va='center')

plt.arrow(matplot_x, matplot_y, 0.2, 0, head_width = 0.05,
↪head_length = 0.05)
ax.annotate(f"{Q[S[x, y]][3]:.2f}", xy=(matplot_x + 0.5,
↪matplot_y), ha='right', va='center')

except Exception as e:
    print(f"Error: {e}")
    print(f"Current x and y: {x}, {y}")

offset = .5
ax.set_xlim(-offset, width - offset)
ax.set_ylim(-offset, height - offset)

ax.hlines(y=np.arange(height+1)- offset, xmin=-offset, xmax=width-offset)
ax.vlines(x=np.arange(width+1) - offset, ymin=-offset, ymax=height-offset)

plt.title(title)

if filename is not None:
    plt.savefig(filename)
    plt.close(fig)
else:
    plt.show()

```

```

[11]: # Defining the plotting function for the policy
def plot_V_Table(
    S: np.ndarray,
    Q: np.ndarray,
    goal_coords: tuple,
    wall_coords: list,
    start_coords: tuple,

```

```

        title: str,
        img_width: int = 6,
        img_height: int = 6,
        filename: str = None
    ):
        height, width = S.shape

        fig = plt.figure(figsize=(img_width, img_height))
        ax = fig.add_subplot(111, aspect='equal')
        for x in range(height):
            for y in range(width):
                matplot_x, matplot_y = array_index_to_matplot_coords(x, y, height)

                # If there is a tuple of (x, y) in the hole_coords list, we color
                ↪the cell gray
                if (x, y) in wall_coords:
                    ax.add_patch(matplotlib.patches.Rectangle((matplot_x - 0.5,
                ↪matplot_y - 0.5), 1, 1, facecolor='gray'))

                # If there is a tuple of (x, y) in the goal_coords list, we color
                ↪the cell yellow
                elif (x, y) in goal_coords:
                    ax.add_patch(matplotlib.patches.Rectangle((matplot_x - 0.5,
                ↪matplot_y - 0.5), 1, 1, facecolor='lightgreen'))

                else:
                    try:
                        # If there is a tuple of (x, y) in the start_coords list,
                ↪we color the cell green
                        if (x, y) in start_coords:
                            ax.add_patch(matplotlib.patches.Rectangle((matplot_x -
                ↪0.5, matplot_y - 0.5), 1, 1, facecolor='yellow'))

                            ax.annotate(f"{max(Q[S[x, y]]):.2f}", xy=(matplot_x,
                ↪matplot_y), ha='center', va='center')

                        # Adding the arrows to the plot
                        if 0 == np.argmax(Q[S[x, y]]):
                            plt.arrow(matplot_x, matplot_y + 0.1, 0, 0.3,
                ↪head_width = 0.05, head_length = 0.05)
                        if 1 == np.argmax(Q[S[x, y]]):
                            plt.arrow(matplot_x, matplot_y - 0.1, 0, -0.3,
                ↪head_width = 0.05, head_length = 0.05)
                        if 2 == np.argmax(Q[S[x, y]]):

```



```

        plt.arrow(matplot_x - 0.1, matplot_y, -0.3, 0,
↪head_width = 0.05, head_length = 0.05)
        if 3 == np.argmax(Q[S[x, y]]):
            plt.arrow(matplot_x + 0.1, matplot_y, 0.3, 0,
↪head_width = 0.05, head_length = 0.05)

    except Exception as e:
        print(f"Error: {e}")
        print(f"Current x and y: {x}, {y}")

offset = .5
ax.set_xlim(-offset, width - offset)
ax.set_ylim(-offset, height - offset)

ax.hlines(y=np.arange(height+1)- offset, xmin=-offset, xmax=width-offset)
ax.vlines(x=np.arange(width+1) - offset, ymin=-offset, ymax=height-offset)

plt.title(title)

if filename is not None:
    plt.savefig(filename)
    plt.close(fig)
else:
    plt.show()

```

### 3 The agent

The agent will be an object created by the class `Agent`. All the uptades of the q table will be done internally in the agent object.

```

[12]: # steps to goal plot
      # add comment about where to add training script

class Agent:
    def __init__(
        self,
        nrow_maze: int,
        ncol_maze: int,
        actions: list = [0, 1, 2, 3],
        rewards: dict = {
            'step': -1,
            'wall': -10,
            'goal': 10,
        },
        gamma: float = 0.9,

```

```

alpha: float = 0.1,
epsilon: float = 0.1,
seed: int = 42,
) -> None:
    """
    Creates an agent for the maze environment.

    Parameters
    -----
    nrow_maze : int
        The number of rows in the maze.
    ncol_maze : int
        The number of columns in the maze.
    actions : list, optional
        A list of actions that the agent can take. The default is [0, 1, 2, 3].
    rewards : dict, optional
        A dictionary of rewards for the agent. The default is {'step': -1, 'wall': -10, 'goal': 10}.
    gamma : float, optional
        The discount factor. The default is 0.9.
    alpha : float, optional
        The learning rate. The default is 0.1.
    epsilon : float, optional
        The exploration rate. The default is 0.1.
    seed : int, optional
        The seed for the random generator. The default is 42.
    """
    self.nrow_maze = nrow_maze
    self.ncol_maze = ncol_maze
    self.rewards = rewards
    self.gamma = gamma
    self.alpha = alpha
    self.epsilon = epsilon
    self.seed = seed
    self.actions = actions

    # By default, the starting index is 0 0
    self.start_state = 0

    # By default, the steps to goal is 0
    self.steps_to_goal = 0

```

```

# Initialize an empty list of steps to goals for all episodes
self.list_steps_to_goals = []

# By default, the goal index is the last index
self.goal_state = nrow_maze * ncol_maze - 1

# Creating the random generator with a fixed seed
self.random_generator = np.random.default_rng(seed)

# Creating the maze; We will denote it internally as S
self.init_S_table()

# Initiating the Q-table
self.init_Q_table()

# Saving the initial past_action and past_state
self.past_action = None
self.past_state = None

# Creating the action name dictionary
self.action_name_dict = {
    0: 'up',
    1: 'down',
    2: 'left',
    3: 'right',
}

# Counter for the number of times our agent has seen the terminal state
self.num_goal_reached = 0

# Counter for each state and how many times the agent visited each
self.state_visit_counter = {}

# Empty dictionary of states visitation paths
self.state_visit_paths = {}

# Empty dictionary of Q values
self.Q_dict = {}

# Placeholder for the current episode of learning
self.current_episode = 0

# Defining the path to intermediate images
self.path_to_images_V = "q-learning-V-Table_Progress"

# Defining the path to intermediate images

```

```

self.path_to_images_Q = "q-learning-Q-Table_Progress"

def increment_state_visit(self, state) -> None:
    """
    Increments the state visit counter for the state.
    """
    if state in self.state_visit_counter:
        self.state_visit_counter[state] += 1
    else:
        self.state_visit_counter[state] = 1

def get_most_recent_action(self) -> str:
    """
    Returns the name of the most recent action.
    """
    return self.action_name_dict[self.past_action]

def init_S_table(self):
    """
    Creates an internal S table where the rows of the table are
    the states and the columns are the actions.
    """
    self.S = np.arange(0, self.nrow_maze * self.ncol_maze).reshape(self.
↪nrow_maze, self.ncol_maze)

def init_Q_table(self):
    """
    Creates an internal q table where the rows of the table are
    the states and the columns are the actions.
    """
    self.Q = np.zeros((self.S.size, len(self.actions)))

def init_reward_dict(self):
    """
    Creates a dictionary where the keys are the states and the values are
↪the rewards for transitioning to that state.
    """
    if self.rewards.get('step') is None:
        raise ValueError("You must specify a reward for taking a step.")
    else:
        self.reward_dict = {s: self.rewards['step'] for s in self.S.
↪flatten()}

    if self.rewards.get('goal') is None:
        raise ValueError("You must specify a reward for reaching the goal
↪state.")
    else:

```

```

        self.reward_dict[self.goal_state] = self.rewards['goal']

    if self.rewards.get('wall') is not None:
        # Setting the reward for the wall states
        for wall_state in self.wall_states:
            self.reward_dict[wall_state] = self.rewards['wall']

def init_maze(self, maze_density: int = None):
    """
    Creates an array of states in a maze environment.

    Returns
    -----
    maze : np.array
        A 2D array of states in a maze environment.
    wall_coords : list
        A 1D array of wall states in the maze environment.
    start_coords: tuple
        A tuple of start coordinates in the maze environment.
    goal_coords: tuple
        A tuple of goal coordinates in the maze environment.
    """
    # If maze_density is None we will create walls in 20% of the maze
    if maze_density is None:
        maze_density = int(self.nrow_maze * self.ncol_maze * 0.2)

    # Creating the wall states
    wall_states = self.random_generator.choice(range(1, self.nrow_maze *
↪self.ncol_maze), size=maze_density, replace=False)

    # Getting the list of wall coordinates
    wall_coords = [self.get_state_coords(s) for s in wall_states]

    # Generating a starting state from the states that are NOT walls
    start_state = self.random_generator.choice(np.setdiff1d(self.S,
↪wall_states))

    # Getting the starting coordinates
    start_coords = self.get_state_coords(start_state)

    # Generating a goal state from the states that are NOT walls and NOT
↪the starting state
    goal_state = self.random_generator.choice(np.setdiff1d(np.
↪setdiff1d(self.S, wall_states), start_state))

    # Getting the goal coordinates
    goal_coords = self.get_state_coords(goal_state)

```

```

    # Saving the wall coordinates, start coordinates, and goal coordinates
    self.wall_coords = wall_coords
    self.start_coords = [start_coords]
    self.goal_coords = [goal_coords]

    # Saving the indexes for the wall, start and goal states
    self.wall_states = wall_states
    self.start_state = start_state
    self.goal_state = goal_state

    # Initiating the reward dictionary
    self.init_reward_dict()

    # Initiating the agent
    self.init_agent()

def get_state_index(self, row: int, col: int) -> int:
    """
    Returns the state index given the state coordinates.

    An inverse function of get_state_coords()
    """
    if (row < 0 or row >= self.nrow_maze or col < 0 or col >= self.
↪ncol_maze):
        return -1
    else:
        return self.S[row][col]

def get_state_coords(self, s) -> tuple:
    """
    Returns the state coordinates given the state index
    """
    s_index = np.where(self.S == s)
    if len(s_index[0]) == 0:
        return -1, -1

    return s_index[0][0], s_index[1][0]

def get_action(
    self,
) -> int:
    """
    Returns a random action from the set of actions

    The actions are:
    0: up

```

```

1: down
2: left
3: right
"""
return self.random_generator.choice(self.actions)

def argmax(self, q_values: np.array):
    """argmax with random tie-breaking
    Args:
        q_values (Numpy array): the array of action values
    Returns:
        action (int): an action with the highest value
    """
    top = float("-inf")
    ties = []

    for i in range(len(q_values)):
        if q_values[i] > top:
            top = q_values[i]
            ties = []

            if q_values[i] == top:
                ties.append(i)

    return self.random_generator.choice(ties)

def get_greedy_action(self, state: int) -> int:
    """
    Returns the greedy action given the current state
    """
    # Getting the q values for the current state
    q_values = self.Q[state]

    # Getting the greedy action
    greedy_action = self.argmax(q_values)

    # Returning the greedy action
    return greedy_action

def get_epsilon_greedy_action(self, state: int) -> int:
    """
    Returns an epsilon greedy action
    """
    if self.random_generator.random() < self.epsilon:
        return self.get_action()
    else:
        return self.get_greedy_action(state)

```

```

# Logging of the agent paths
def log_agent_move(self, state: int):
    """
    Logs the agent's move
    """
    if self.state_visit_paths.get(self.current_episode, None) is None:
        self.state_visit_paths[self.current_episode] = [state]
    else:
        self.state_visit_paths[self.current_episode].append(state)

def update_Q_table(self, new_state: int):
    """
    Function that applies the RL update function
    """
    # Getting the next_state's reward
    reward = self.reward_dict[new_state]

    # Saving the current Q value
    current_Q = self.Q[self.past_state][self.past_action]

    # If the new state is the terminal state or the wall state, then the
    ↪ max_Q is 0
    max_Q = 0

    # Else we get the max Q value for the new state
    if new_state != self.goal_state:
        new_state_Q_values = self.Q[new_state]

        # Getting the max Q value
        max_Q = np.max(new_state_Q_values)

    # Updating inplace the Q value
    self.Q[self.past_state][self.past_action] = current_Q + self.alpha *
    ↪ (reward + self.gamma * max_Q - current_Q)

def get_next_state(self, s: int, action: int) -> int:
    """
    Given the current state and the current action, returns the next state
    ↪ index
    """
    # Getting the state coordinates
    s_row, s_col = self.get_state_coords(s)

```



```

    # Setting the boolean indicating that we have reached the terminal state
    reached_terminal = False

    # Getting the next state
    next_state = -1
    if action == 0:
        next_state = self.get_state_index(s_row - 1, s_col)
    elif action == 1:
        next_state = self.get_state_index(s_row + 1, s_col)
    elif action == 2:
        next_state = self.get_state_index(s_row, s_col - 1)
    elif action == 3:
        next_state = self.get_state_index(s_row, s_col + 1)

    # If next_state is a wall or the agent is out of bounds, we will stay
    ↪ in the same state
    if (next_state == -1) or (next_state in self.wall_states):
        return s, reached_terminal

    # Incrementing the number of times we have visited the next state
    self.increment_state_visit(next_state)

    # If next_state is the goal state, we will return to the starting state
    if next_state == self.goal_state:
        # Incrementing the number of times our agent has reached the goal
        ↪ state
        self.num_goal_reached += 1
        reached_terminal = True

    # Returning the next state
    return next_state, reached_terminal

def init_agent(self):
    """
    We will set the past state and past action as the starting state and
    ↪ action
    """
    # Setting the previous state as the starting state
    self.past_state = self.start_state
    self.past_action = self.get_epsilon_greedy_action(self.past_state)
    self.num_goal_reached = 0

def save_steps_to_goal_for_current_episode(self):
    """
    Write your own code here to save the steps to goal for current episode
    """

```

```

        #print(f"Episode {self.current_episode}: Steps to goal = {self.
↪steps_to_goal}")
        #In list_steps_to_goals list, steps_to_goal value is appended to have
↪list of goals for all episodes
        self.list_steps_to_goals.append(self.steps_to_goal)

        return self.list_steps_to_goals

def terminal_step(self, new_state: int):
    """
    Updates the agent one last time and resets the agent to the starting
↪position
    """
    # Updating the Q table
    self.update_Q_table(new_state)

    # Resetting the agent
    self.past_state = self.start_state
    self.past_action = self.get_epsilon_greedy_action(self.past_state)

    # save the steps taken to reach the goal for this episode, and reset
↪steps to goal to 0 for next episode
    self.save_steps_to_goal_for_current_episode()
    self.steps_to_goal = 0

    # Incrementing the number of episodes
    self.current_episode += 1

def move_agent(self):
    """
    The function that moves the agent to the next state
    """
    # Getting the next state
    next_state, reached_terminal = self.get_next_state(self.past_state,
↪self.past_action)

    # Adding the next state to the path
    self.log_agent_move(next_state)

    # Updating the Q table
    if not reached_terminal:
        # Checking if the past_state is the same as the next_state; If that
↪is true, it means our agent hit a wall

```

```

        # or went out of bounds
        if self.past_state != next_state:
            self.update_Q_table(next_state)

        # Setting the past_state as the next_state
        self.past_state = next_state

        # Getting the next action
        self.past_action = self.get_epsilon_greedy_action(self.past_state)

        # add 1 step
        self.steps_to_goal += 1

    else:
        # add 1 step
        self.steps_to_goal += 1
        self.terminal_step(next_state)

def train_episodes(self, num_episodes: int):
    """
    Function that trains the agent for one episode
    """
    # Calculating the episode number to end the training
    end_episode = self.current_episode + num_episodes - 1

    # Moving the agent until we reach the goal state
    while self.current_episode != end_episode:

        # print(f"Episode {self.current_episode}")

        self.move_agent()

def create_policy(self):
    """
    Creates a policy dictionary where the key is the state and the value is
    ↪ the action
    based on the Q table
    """
    # Creating the policy dictionary
    self.policy = {}

    # Looping through the states
    for state in range(self.S.size):
        # Getting the greedy action
        greedy_action = self.argmax(self.Q[state])

```

```

        # Adding the state and action to the policy dictionary
        self.policy[state] = greedy_action

def create_optimal_policy_path(self):
    """
    Creates the path of the optimal policy, starting from the starting state
    """
    # Creating the policy path
    self.optimal_policy_path = [self.start_state]
    self.optimal_policy_path_coords = [self.get_state_coords(self.
↪start_state)]

    # Getting the current state
    current_state = self.start_state

    # Looping through the states
    while current_state != self.goal_state:
        # Getting the next state
        next_state, _ = self.get_next_state(current_state, self.
↪policy[current_state])

        # Adding the next state to the path
        self.optimal_policy_path.append(next_state)
        self.optimal_policy_path_coords.append(self.
↪get_state_coords(next_state))

        # Setting the current state as the next state
        current_state = next_state

```

## 4 Agent learning

First, let us create the agent object and initiate the maze environment for it.

### 5 Q1. Design your own grid world problem: you can design arbitrary grid world problem but should be bigger than 4\*3; the grid world must include obstacles or walls

- Grid world Configuration:
  - Size: 5x4 maze
  - Walls: 2 walls placed within the maze

### 6 Q2: Design your states, actions, and rewards.

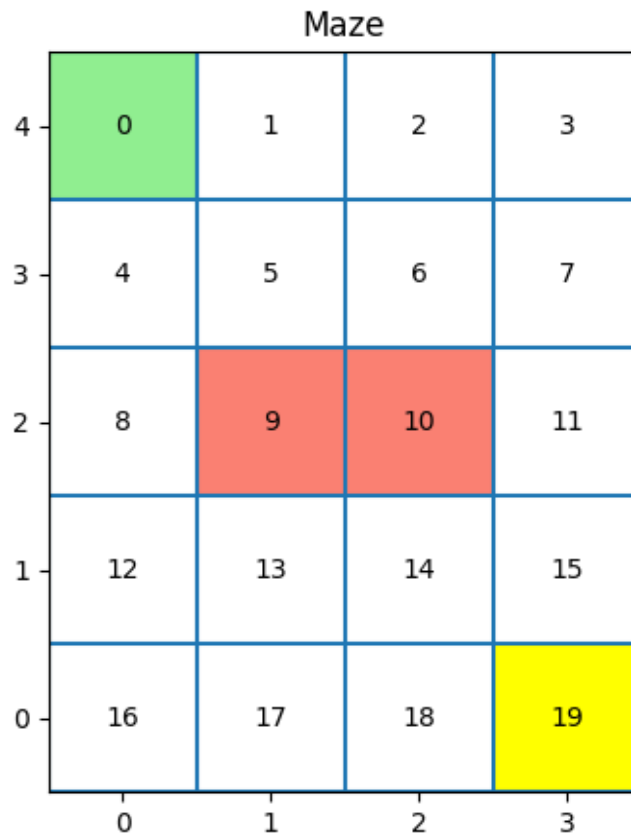
- States:

- Initial: on coords (3,0) which is in yellow color
- Goal: in (0,4) which is in green color
- Actions: There are four actions at every state: up, down, right, left.
- Rewards:
  - Step: 0
  - Wall: -5
  - Goal: 10

```
[13]: agent = Agent(
    nrow_maze=5,
    ncol_maze=4,
    seed=1,
    rewards={'step': 0, 'wall': -5, 'goal': 10}
)

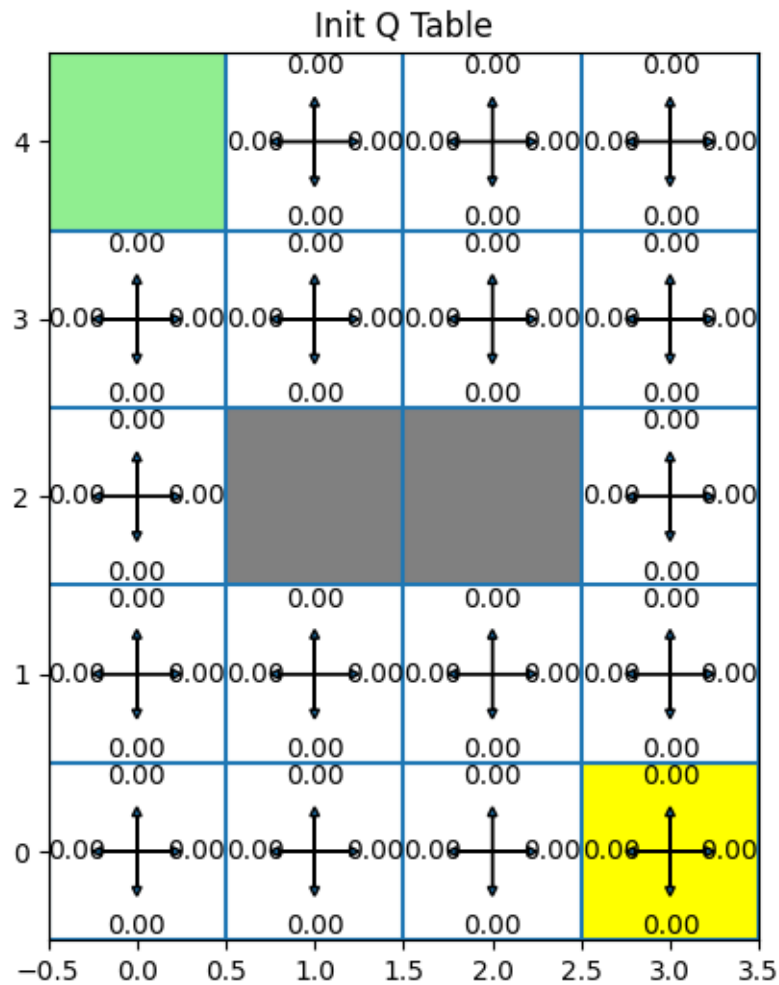
# Initiating the maze
agent.init_maze(maze_density=2)

# Plotting the maze
plot_matrix(agent.S, goal_coords=agent.goal_coords, hole_coords=agent.
    ↪ wall_coords, start_coords=agent.start_coords, title="Maze")
```



Display the initial Q table

```
[14]: plot_Q_Table(  
    S=agent.S,  
    Q=agent.Q,  
    goal_coords=agent.goal_coords,  
    wall_coords=agent.wall_coords,  
    start_coords=agent.start_coords,  
    title="Init Q Table"  
)
```



## 7 Q3. Implement your Q-learning algorithm.

```
[15]: # Training the agent for one episode
agent.train_episodes(num_episodes=100)
```

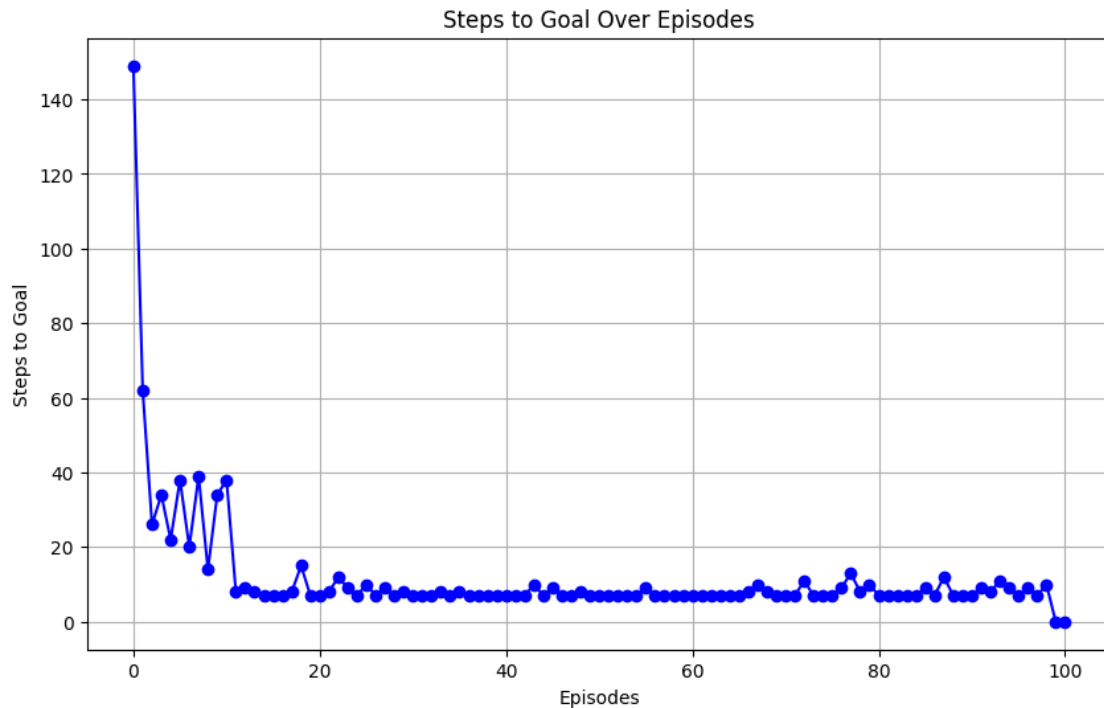
## 8 Q4. Show your goal searching process with step-to-go curve, and learned Q-table (good to show some intermediate Q-tables).

#step-to-go curve

```
[24]: steps_to_goal_list_for_all = agent.save_steps_to_goal_for_current_episode()
```

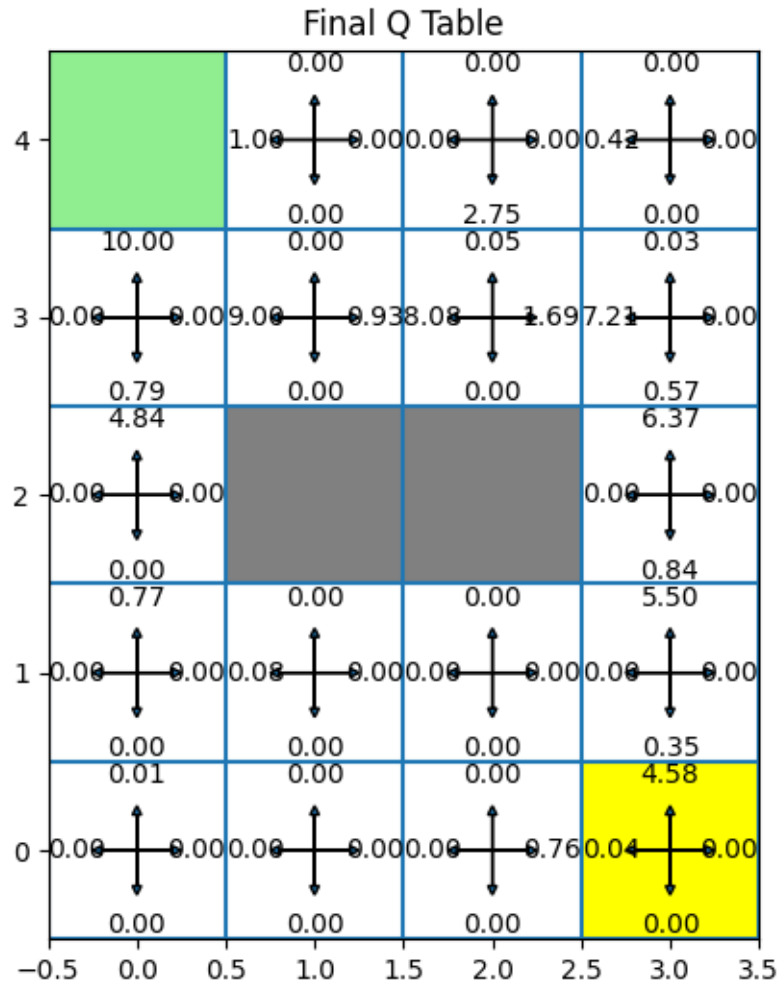
```
[26]: import matplotlib.pyplot as plt

# Plotting the graph
plt.figure(figsize=(10, 6))
plt.plot(steps_to_goal_list_for_all, marker='o', linestyle='-', color='b')
plt.title('Steps to Goal Curve')
plt.xlabel('Episodes')
plt.ylabel('Steps to Goal')
plt.grid(True)
plt.show()
```



#learned Q-table

```
[27]: plot_Q_Table(
    S=agent.S,
    Q=agent.Q,
    goal_coords=agent.goal_coords,
    wall_coords=agent.wall_coords,
    start_coords=agent.start_coords,
    title="Final Q Table"
)
```



#Intermediate Q tables

```
[34]: def q_learning_for_intermediate_episodes(num_episodes):

    agent.train_episodes(num_episodes)
    plot_Q_Table(
        S=agent.S,
```



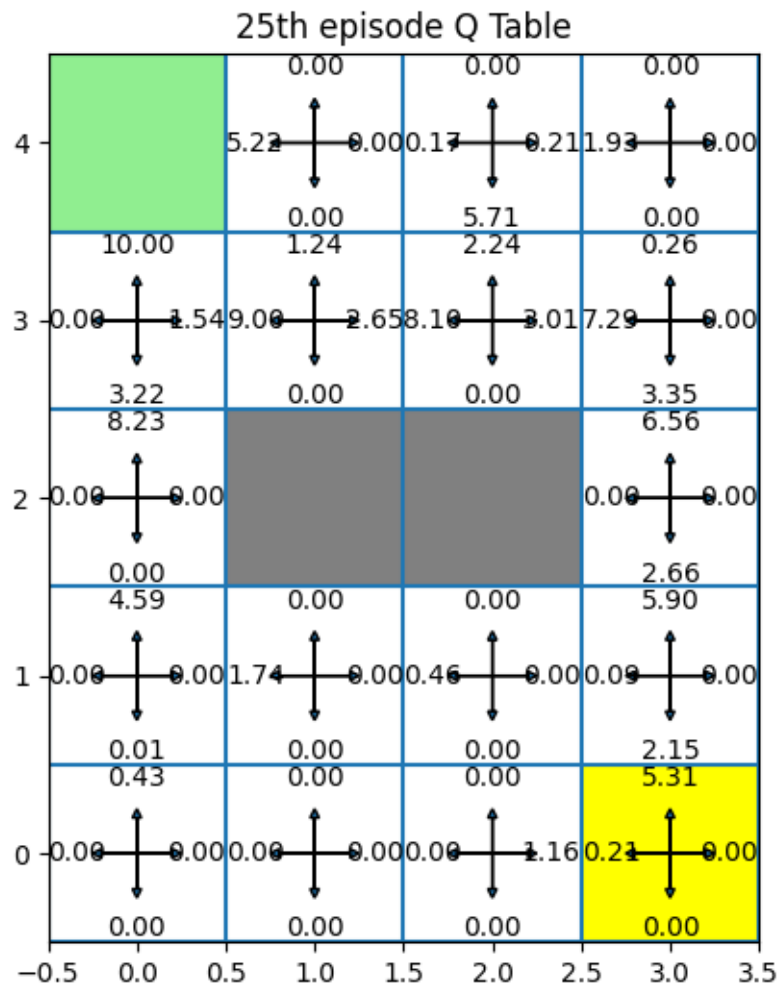
```

Q=agent.Q,
goal_coords=agent.goal_coords,
wall_coords=agent.wall_coords,
start_coords=agent.start_coords,
title=f"{num_episodes}th episode Q Table"
)

```

##Q learning for 25th Episode

```
[35]: q_learning_for_intermediate_episodes(25)
```



##Q learning for 50th Episode

```
[36]: q_learning_for_intermediate_episodes(50)
```

50th episode Q Table

